

BIAS-VARIANCE TRADEOFF

Bias

Error from an **overly simple** model. High bias misses the real pattern even on the training set: it **underfits**.



Variance

Error from **sensitivity to the training set**. High variance fits the noise in one sample and does not generalize: it **overfits**.



Low variance

High variance


Low bias



High bias



OVERFITTING AND UNDERFITTING



Underfitting

Overfitting

Underfitting

A boundary too simple for the data: it misses points on both sides. High **bias**.

Overfitting

A boundary that hugs every training point, noise included. High **variance**.

Good fit

A boundary that follows the pattern and ignores the noise: the **sweet spot**.

WAYS TO MITIGATE OVERFITTING AND UNDERFITTING

If the model overfits

- ▶ **More data**, or **augmentation** to synthesize more.
- ▶ **Regularization**: weight decay, dropout.
- ▶ **Early stopping** at the validation minimum.
- ▶ **Reduce capacity**: fewer layers or parameters.
- ▶ Start from **pretrained weights** instead of scratch.

If the model underfits

- ▶ **Train longer**, or raise the learning rate.
- ▶ **Add capacity**: more layers or parameters.
- ▶ **Add features**, or a richer input representation.
- ▶ **Reduce regularization**: less weight decay or dropout.
- ▶ Check the loss and labels: a bug can look like underfitting.

TERMINOLOGY

		Prediction	
		0	1
Ground truth	0	TN m_{00}	FP m_{01}
	1	FN m_{10}	TP m_{11}

↑ missed disease

Be careful: some packages transpose this layout.

- ▶ **True positive** (m_{11}): an example of class 1 predicted as class 1.
- ▶ **False positive** (m_{01}): class 0 predicted as class 1. **Type I error.**
- ▶ **True negative** (m_{00}): class 0 predicted as class 0.
- ▶ **False negative** (m_{10}): class 1 predicted as class 0. **Type II error.**

Total instances $m = m_{00} + m_{01} + m_{10} + m_{11}$, and the **error rate** is $(m_{01} + m_{10})/m$.

Error rate hides imbalance. If 10% of examples are class 1 and 90% are class 0, calling *everything* class 0 already scores a 10% error rate. On a disease with 5% prevalence, always predicting “healthy” reaches **95% accuracy** and catches **zero** patients: every error it makes is an m_{10} , the one cell that matters clinically.

COMMON MEASURES



What machine learning calls them

- ▶ Accuracy = $\frac{TP + TN}{TP + FP + FN + TN}$
- ▶ Precision = $\frac{TP}{TP + FP}$ of the declared positives
- ▶ Recall = $\frac{TP}{TP + FN}$ of the actual positives
- ▶ $F_1 = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$


What medicine calls them

- ▶ Sensitivity = $\frac{TP}{TP + FN}$ the same as recall
- ▶ Specificity = $\frac{TN}{FP + TN}$ of the actual negatives
- ▶ False positive rate = $\frac{FP}{FP + TN}$
- ▶ PPV = precision, NPV = $\frac{TN}{TN + FN}$

CLASSIFICATION: ROC AND AUC



A single number for the **whole curve**: higher AUC means better separation between classes.



ROC curves

Sweep the threshold from 1 to 0 and plot sensitivity against 1—specificity: threshold-free, good for comparing models.

Best-performing thresholds

AUC does not deploy. You must pick a point: **A** for high specificity (few false alarms), **C** for high sensitivity (screening), **B** as a balanced compromise.